

KRIA board Ubuntu LTS 22.04 Install

Vincent Conus*

2023-12-22

Contents

1	Preparing and booting a Ubuntu 22.04 media	1
2	Network and admin setups	2
2.1	Static IP address	2
2.2	[facultative] root password	3
2.3	Adding Xilinx specific repositories	3
2.4	[facultative] Installing Docker	3
2.5	[recommended] Purging snap	4
2.6	[facultative] Other unused heavy packages	5
2.7	[facultative] Slow boot services to disable	5
2.8	[facultative] Adding a swap partition	6
2.9	[OLD] Enabling remoteproc with Device-Tree patching	6
2.10	[NEW] Enabling remoteproc for RPMsg in userspace with device-tree patch	8
2.10.1	Patching the device tree for RPMsg in userspace	8
2.10.2	Checking the patching	8
2.11	[DEPRECATED] Proxy and DNS	8
2.12	[DEPRECATED] Jupyter notebook setup	8
2.13	[TODO] Using a PetaLinux kernel in Ubuntu	9
3	RPMsg standalone evaluation	10
3.1	[TODO] RPMsg Cortex R5F demonstration firmware	10
3.2	RPMsg echo_test software	10

1 Preparing and booting a Ubuntu 22.04 media

An official Ubuntu image exists and is provided by Xilinx, allowing the OS installation to be quick and straight-forward. Ubuntu is a common and easy to use distribution. Furthermore, it allows to install ROS2 as a package, which is most convenient and will be done later in this guide.

Once the image has been downloaded at Canonical's page we can flash it onto the SD card, with the following instructions.

DANGER: The next part involve the dd command writing on disks!!! As always with the dd command, thou have to be **VERY** careful on what arguments thou give. Selecting the wrong disk will result on the destruction of thy data !! If you are unsure of what to do, seek assistance !

With the image available on thy machine and a SD card visible as /dev/sda device¹ one can simply run the dd command as follow to write the image to a previously formatted drive (here /dev/sda):

*vincent.conus@protonmail.com

¹Again, it is critical to be 100\the correct device!

```

1 unxz iot-limerick-kria-classic-desktop-2204-20240304-165.img.xz
2 sudo dd if=iot-limerick-kria-classic-desktop-2204-20240304-165.img \
3 of=/dev/sda status=progress bs=8M && sync

```

Once the SD card is flashed and put back in the board, the micro-USB cable can be connected from the PC to the board. It is then possible to connect to the board in serial with an appropriate tool, for example `picocom`, as in the following example (the serial port that "appeared" was the `/dev/ttyUSB1` in this case², and the 115200 bit-rate is the default value for the board):

```

1 sudo picocom /dev/ttyUSB1 -b 115200

```

In my case, I am using Emacs's `serial-term`:

```

1 M-x serial-term RET /dev/ttyUSB1 RET 115200 RET

```

The default username / password pair for the very first boot is `ubuntu` and `ubuntu`. You will then be prompted to enter a new password.

Once logged in, it is typically easier and more convenient to connect the board using SSH. When the board is connected to the network, it is possible to know its IP address with the `IP` command; then it is possible to connect to the board with `ssh`, as follow (example, with the first command to be run on the board and the second one on the host PC, both without the first placeholder hostnames):

```

1 kria# ip addr
2
3 host# ssh ubuntu@192.168.4.11

```

2 Network and admin setups

This section presents a variety of extra convenience configurations that can be used when setting-up the Kria board.

2.1 Static IP address

A static IP can be set by writing the following configuration into your `netplan` configuration file³.

The name of the files might vary:

```

1 sudo chmod 0600 /etc/netplan/50-cloud-init.yaml
2 sudo nano /etc/netplan/50-cloud-init.yaml

```

You can then set the wanted IP as follow⁴:

```

1 network:
2   renderer: NetworkManager
3   version: 2
4   ethernets:
5     eth0:
6       dhcp4: false

```

²If two boards are plugged in for serial, the second one will be `/dev/ttyUSB5`, then `USB9` and so on.

³The `chmod` command is used to update the permissions and silence some warnings

⁴For the routing part, it is key to have the `to` with a `'-'` in front of it; and then the `via` without, but aligned with the `t`.

```

7      addresses:
8        - 192.168.11.107/24
9      routes:
10       - to: default
11         via: 192.168.11.1
12      nameservers:
13        addresses: [192.168.11.1]

```

Finally, the change in settings can be applied as follow:

```

1  sudo netplan apply

```

2.2 [facultative] root password

WARNING: Depending on your use-case, the setup presented in this subsection can be a critical security breach as it remove the need for a root password to access the admin functions of the board's Linux. When in doubt, do not apply this configuration!!

If you board does not hold important data and is available to you only, for test or development, it might be convenient for the sudo tool to not ask for the password all the time. This change can be done by editing the sudoers file, and adding the parameter NOPASSWD at the sudo line:

```

1  sudo visudo
2
3  %sudo    ALL=(ALL:ALL) NOPASSWD: ALL

```

Again, this is merely a convenience setup for devices staying at you desk. If the board is meant to be used in any kind of production setup, a password should be set for making administration tasks.

With all of these settings, you should be able to update the software of your board without any issues:

```

1  sudo apt update
2  sudo apt dist-upgrade
3  sudo reboot now

```

2.3 Adding Xilinx specific repositories

The following commands will add PPA repositories that are specific for Xilinx boards using Ubuntu. It is then possible to update the package list and eventually upgrade to some new packages.

```

1  sudo add-apt-repository ppa:ubuntu-xilinx/updates
2  sudo add-apt-repository ppa:xilinx-apps/ppa
3  sudo apt update
4  sudo apt upgrade

```

2.4 [facultative] Installing Docker

It is possible to have a version of Docker installed simply by using the available repository, but since we are on Ubuntu, a PPA is available from Docker in order to have the most up-to-date version.

Following the official documentation, the following steps can be taken to install the latest version of Docker on a Ubuntu system. The last command is meant to test the install. If everything went smoothly, you should see something similar to what is presented in the figure 1 below, after the commands:

```

1  sudo apt-get update
2  sudo apt-get install ca-certificates curl
3  sudo install -m 0755 -d /etc/apt/keyrings
4  curl -fsSL https://download.docker.com/linux/ubuntu/gpg | \
5      sudo gpg --dearmor -o /etc/apt/keyrings/docker.gpg
6
7  sudo chmod a+r /etc/apt/keyrings/docker.gpg
8
9  echo \
10     "deb [arch="$(dpkg --print-architecture)" \
11     signed-by=/etc/apt/keyrings/docker.gpg] \
12     https://download.docker.com/linux/ubuntu \
13     "$(. /etc/os-release && \
14     echo "$VERSION_CODENAME")" stable" | \
15     sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
16
17  sudo apt-get update
18  sudo apt-get install docker-ce docker-ce-cli \
19     containerd.io docker-buildx-plugin docker-compose-plugin
20  sudo usermod -aG docker $USER
21  newgrp docker
22
23  docker run hello-world

```

```

ubuntu@kria:~$ docker run hello-world
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
70f5ac315c5a: Pull complete
Digest: sha256:fc6cf906cbfa013e80938cdf0bb199fbd6b86d6e3e013783e5a766f50f5dbce0
Status: Downloaded newer image for hello-world:latest

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
 1. The Docker client contacted the Docker daemon.
 2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
    (arm64v8)
 3. The Docker daemon created a new container from that image which runs the
    executable that produces the output you are currently reading.
 4. The Docker daemon streamed that output to the Docker client, which sent it
    to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
 $ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
 https://hub.docker.com/

For more examples and ideas, visit:
 https://docs.docker.com/get-started/

```

Figure 1: The return of a successful run of the hello world test Docker container.

2.5 [recommended] Purging snap

As the desktop-specific software are not used at all in the case of our project, there are some packages that can be purged in order for the system to become more lightweight.

In particular, the main issue with Ubuntu systems is the forced integration of Snap packages. Here are the commands to use in order to remove all of that. These steps take a lot of time and need to be executed in that

specific order⁵, but the system fan runs sensibly slower without all of this stuff.

Note that this all process is rather slow and can take up to 30min to complete.

```
1 sudo systemctl disable snapd.service
2 sudo systemctl disable snapd.socket
3 sudo systemctl disable snapd.seeded.service
4
5 sudo snap list #show installed package, remove them all:
6 sudo snap remove --purge firefox
7 sudo snap remove --purge gnome-3-38-2004
8 sudo snap remove --purge gnome-42-2204
9 sudo snap remove --purge gtk-common-themes
10 sudo snap remove --purge snapd-desktop-integration
11 sudo snap remove --purge snap-store
12 sudo snap remove --purge bare
13 sudo snap remove --purge core20
14 sudo snap remove --purge core22
15 sudo snap remove --purge snapd
16 sudo snap list # check that everything is uninstalled
17
18 sudo rm -rf /var/cache/snapd/
19 sudo rm -rf ~/snap
20 sudo apt autoremove --purge snapd
21
22 # check once more that there is no more snap on the system
23 systemctl list-units | grep snapd
```

2.6 [facultative] Other unused heavy packages

Some other pieces of software can safely be removed since the desktop is not to be used.

```
1 sudo apt-get autoremove --purge yaru-theme-icon \
2     fonts-noto-cjk yaru-theme-gtk vim-runtime \
3     ubuntu-wallpapers-jammy humanity-icon-theme
4
5 sudo apt-get autoclean
6 sudo reboot now
```

2.7 [facultative] Slow boot services to disable

These packages (in particular the first one) are taking up a LOT of time at boot while providing no benefits⁶.

It is possible to disable them as follow:

```
1 sudo systemctl disable systemd-networkd-wait-online.service
2 sudo systemctl disable NetworkManager-wait-online.service
3 sudo systemctl disable cups.service
4 sudo systemctl disable docker.service
5 sudo systemctl disable containerd.service
6 sudo systemctl disable cloud-init-local.service
```

Additional, potentially unused services can be found using the very handy command:

⁵The snap packages depends on each others. Dependencies cannot be remove before the package(s) that depends on them, thus the specific delete order.

⁶The CUPS and Docker services will be activated when used instead of during boot time.

```
1 sudo systemd-analyze blame
```

2.8 [facultative] Adding a swap partition

This part is very optional, in particular as it might slow down a bit the boot time of the board (~2s), however it might become handy to have swap memory available to avoid system failure under heavy use.

This whole procedure must be done externally, with the board system SD card mounted on a host PC as an external volume. As it is highly platform dependant, I will not give a detailed explanation on how to do it, yet here are the key points that should be done:

- Shutdown the Kria board, take out the SD card and put it in a host machine.
- Make sure the disk is visible.
- Make sure all volumes are **unmounted**.
- Resize the main root partition (**not** the boot) so a space the size of the wanted swap is free **after** the partition. You'd want something around 1GB.
- In the empty space, create a new partition, which type is "linux swap".
- Find and take note of the UUID of the new partition. This is useful hereafter.
- `sync`
- Un-mount everything, eject SD card.
- Put the SD card back in the Kria.
- Boot back to Ubuntu.

Going back on the Kria board Ubuntu after boot, the `/etc/fstab` file can be updated as follow, modulo your actual UUID for the newly created partition, to enable swap at boot time.

```
1 sudo -s
2 echo "UUID=8b13ed05-a91d-4x50-a44a-e654a0c67a2c none swap sw 0 0" >> /etc/fstab
3 reboot now
```

2.9 [OLD] Enabling remoteproc with Device-Tree patching

One of the advantage of this Kria board, as cited previously, is the presence of multiple types of core (APU, MCU, FPGA) on the same chip.

The part in focus in this guide is the usage of both the APU, running a Linux distribution and ROS2; and the MCU, running FreeRTOS and micro-ROS. Online available guides⁷ · ⁸ also provide information on how to deploy these types of systems and enabling remoteproc for the Kria board, but this guide will show a step-by-step, tried process to have a heterogeneous system up and running.

The communication between both side is meant to be done using shared memory, but some extra setup is required in order to be running the real-time firmware, in particular for deploying micro-ROS on it.

As a first step in that direction, this section of the report will present how to setup and use as an example firmware that utilizes the remoteproc device in Linux in order to access shared memory and communicate with the real-time firmware using the RPMsg system.

The communication system and interaction from the Linux side towards the real-time capable core is not enabled by default within the Ubuntu image provided by Xilinx.

⁷A slideshow (JP) from Fixstar employees presents how to use the device tree to enable the communication between the cores.

⁸A blog post (JP) shows all major steps on how to enable the remoteproc.

In that regard, some modification of the device tree overlay (DTO) is required in order to have the remoteproc system starting.

Firstly, we need to get the original firmware device tree, converted into a readable format (DTS):

```
1 sudo dtc /sys/firmware/fdt 2> /dev/null > system.dts
```

Then, a custom-made patch file can be downloaded and applied. This file is available at the URL visible in the command below.

```
1 wget  
2 ↪ https://gitlab.com/sunoc/xilinx-kria-kv260-documentation/-/raw/806dcaa16ed9f2571539fb98e9f027226a485a6a/src/system.patch  
3 patch system.dts < system.patch
```

As for the board to be able to reserve the correct amount of memory with the new settings, some cma kernel configuration is needed⁹:

```
1 sudo nano /etc/default/flash-kernel  
2  
3 LINUX_KERNEL_CMDLINE="quiet splash cma=512M cpuidle.off=1"  
4 LINUX_KERNEL_CMDLINE_DEFAULTS=""  
5 sudo flash-kernel
```

Now the DTS file has been modified, one can regenerate the binary and place it on the /boot partition and reboot the board:

```
1 dtc -I dts -O dtb system.dts -o user-override.dtb  
2 sudo mv user-override.dtb /boot/firmware/  
3 sudo reboot now
```

After rebooting, you can check the content of the remoteproc system directory, and a remoteproc0 device should be visible, as follow:

```
1 ls /sys/class/remoteproc/  
2 # remoteproc0
```

If it is the case, it means that the patch was successful and that the remote processor is ready to be used!

⁹The overlapping memory will not prevent the board to boot, but it disables the PWM for the CPU fan, which will then run at full speed, making noise.

2.10 [NEW] Enabling remoteproc for RPMsg in userspace with device-tree patch

2.10.1 Patching the device tree for RPMsg in userspace

The system_uio.patch

```
1 sudo dtc /sys/firmware/fdt 2> /dev/null > system.dts
2
3 wget https://gitlab.com/sunoc/xilinx-kria-kv260-documentation/-/raw/main/src/system_uio.patch
4
5 patch system.dts < system_uio.patch
6
7 dtc -I dts -O dtb system.dts -o user-override.dtb
8 sudo mv user-override.dtb /boot/firmware/
9 sudo reboot now
```

2.10.2 Checking the patching

If everything went correctly, on reboot and re-logging in the device, you should see an output to the following command:

```
1 ls /sys/class/remoteproc/
```

2.11 [DEPRECATED] Proxy and DNS

WARNING: This setup is not needed anymore when using Honda-sensei's lab wired network. Adding it will cause the DNS to FAIL! This section is kept as a reference.

An issue that can occur when connecting the board to the internet is the conflicting situation with the university proxy. Indeed, as the network at Nanzan University requires to go through a proxy, some DNS errors appeared.

In that case, it might become needed to setup the proxy for the school.

This can be done as follow, by exporting a https base proxy configuration containing you AXIA credentials (this is specific to Nanzan University IT system), then by consolidating the configuration for other types of connections in the bashrc:

```
1 export https_proxy="http://<AXIA_username>:\
2     <AXIA_psw>@proxy.ic.nanzan-u.ac.jp:8080"
3
4 echo "export http_proxy=\"\$https_proxy\"" >> ~/.bashrc
5 echo "export https_proxy=\"\$https_proxy\"" >> ~/.bashrc
6 echo "export ftp_proxy=\"\$https_proxy\"" >> ~/.bashrc
7 echo "export no_proxy=\"localhost, 127.0.0.1, ::1\"" >> ~/.bashrc
```

Eventually the board can be rebooted in order for the setup to get applied cleanly.

2.12 [DEPRECATED] Jupyter notebook setup

Here are some instruction on how to install and setup Jupyter on a KRIA board, accessing it remotely and using it for making data analysis.

The following commands will set the required packages and install Jupyter itself¹⁰:

¹⁰Alongside other packages useful for data analysis, such as pandas or numpy.


```

1 sudo apt-get update && sudo apt-get install python3 python3-pip python3-venv python3-virtualenv
2
3 virtualenv myjupyter
4 source ./myjupyter/bin/activate
5 python3 -m pip install jupyter pandas numpy matplotlib scipy
6
7 sudo reboot now

```

Then in a terminal on your host machine (not on the KRIA board), you can run the following command¹¹ to bind local ports:

```

1 ssh -L 8888:localhost:8888 ubuntu@192.168.11.107

```

Then on the opened SSH shell to the KRIA board:

```

1 source ./myjupyter/bin/activate
2 jupyter notebook

```

From there, it is possible to use the displayed URL (something that looks like <http://localhost:8888/tree?token>) to access the remote Notebook system from a local web browser. It is possible to do so with localhost since we have the ssh port map connection going on.

Eventually creating Notebooks and stuff, it is possible to obtain a situation like shown in the figure 2 below.



Figure 2: A test Jupyter Notebook for CSV data analysis.

2.13 [TODO] Using a PetaLinux kernel in Ubuntu

TBD

¹¹In this example, the full username@IP is used, but a .ssh/config is also usable.

3 RPMsg standalone evaluation

3.1 [TODO] RPMsg Cortex R5F demonstration firmware

3.2 RPMsg echo_test software

In order to test the deployment of the firmware on the R5F side, and in particular to test the RPMsg function, we need some program on the Linux side of the KRIA board to "talk" with the real-time side.

Some source is provided by Xilinx to build a demonstration software that does this purpose: specifically interact with the demonstration firmware.

Here are the steps required to obtain the sources, and build the program.

As a reminder, this is meant to be done on the Linux running on the KRIA board, NOT on your host machine

!

```
1 git clone https://github.com/Xilinx/meta-openamp.git
2 cd meta-openamp
3 git checkout xlnx-rel-v2022.2
4 cd ./recipes-openamp/rpmsg-examples/rpmsg-echo-test
5 make
6 sudo ln -s $(pwd)/echo_test /usr/bin/
```

Once this is done, it is possible to run the test program from the KRIA board's Ubuntu by running the `echo_test` command.