

# Setting up and using Xilinx KRIA KV260 and KR260

南山大学

2025-03-04

Vincent Conus - Source available at [GitLab](#) 



# Contents

<b>1</b>	<b>Introduction &amp; motivation</b>	<b>3</b>
<b>2</b>	<b>Boot firmware update</b>	<b>4</b>
2.1	Getting the new firmware . . . . .	4
2.2	Reaching the board recovery tool . . . . .	4
2.3	Updating the boot firmware . . . . .	4
<b>3</b>	<b>Installing Ubuntu LTS 22.04</b>	<b>7</b>
3.1	Preparing and booting a Ubuntu 22.04 media . . . . .	7
3.2	Network and admin setups . . . . .	8
3.2.1	Static IP address . . . . .	8
3.2.2	[facultative] root password . . . . .	9
3.2.3	Adding Xilinx specific repositories . . . . .	9
3.2.4	[facultative] Installing Docker . . . . .	9
3.2.5	[recommended] Purging snap . . . . .	10
3.2.6	[facultative] Other unused heavy packages . . . . .	11
3.2.7	[facultative] Slow boot services to disable . . . . .	11
3.2.8	[facultative] Adding a swap partition . . . . .	11
3.2.9	[OLD] Enabling remoteproc with Device-Tree patching . . . . .	12
3.2.10	[NEW] Enabling remoteproc for RPMsg in userspace with device-tree patch . . . . .	14
3.2.11	[DEPRECATED] Proxy and DNS . . . . .	14
3.2.12	[DEPRECATED] Jupyter notebook setup . . . . .	15
3.2.13	[TODO] Using a PetaLinux kernel in Ubuntu . . . . .	16
3.3	RPMsg standalone evaluation . . . . .	16
3.3.1	[TODO] RPMsg Cortex R5F demonstration firmware . . . . .	16
3.3.2	RPMsg echo_test software . . . . .	16
<b>4</b>	<b>Building micro-ROS as a static library</b>	<b>16</b>
4.1	Initial setup . . . . .	16
4.2	Building the static library . . . . .	17
<b>5</b>	<b>ROS2 host system setup</b>	<b>18</b>
5.1	On the host Linux ("bare-metal") . . . . .	18
5.2	In a container (Docker) . . . . .	20
<b>6</b>	<b>Micro-ROS Client</b>	<b>WORK_IN_PROGRESS 22</b>
6.1	Libraries . . . . .	22
6.2	Prerequisite: KRIA board setup . . . . .	22
6.3	Building the firmware . . . . .	22
6.4	Deployment . . . . .	22
6.5	Starting the Agent . . . . .	23
6.6	Starting the Talker . . . . .	23
6.7	Other useful functions and notes . . . . .	23
6.7.1	"Stop and Reboot" button . . . . .	23
6.7.2	Building the micro-ROS static library . . . . .	23
6.7.3	RPMsg in userspace . . . . .	24
<b>7</b>	<b>micro-ROS XRCE-DDS Agent</b>	<b>WORK_IN_PROGRESS 24</b>
7.1	Building a XRCE-DDS agent in a Docker . . . . .	24
<b>8</b>	<b>Talker / ping-pong node</b>	<b>WORK_IN_PROGRESS 25</b>
<b>9</b>	<b>Monitoring and performances evaluation</b>	<b>WORK_IN_PROGRESS 26</b>

<b>10 Conclusion &amp; future</b>	<b>WORK_IN_PROGRESS</b>	<b>27</b>
<b>11 Publications</b>	<b>WORK_IN_PROGRESS</b>	<b>28</b>
<b>A DTO patch</b>		<b>29</b>
<b>B Custom tool-chain CMake settings</b>		<b>32</b>
<b>C Custom Colcon meta settings</b>		<b>33</b>
<b>D Firmware time functions</b>		<b>35</b>
D.1 main . . . . .		35
D.2 header file . . . . .		36
<b>E Firmware memory allocation functions</b>		<b>37</b>
E.1 main . . . . .		37
E.2 header file . . . . .		38
<b>F [deprecated] Installing Linux (PetaLinux option)</b>		<b>38</b>
F.1 General installation . . . . .		39
F.2 First login . . . . .		41
F.3 Testing openamp "echo-test" . . . . .		41
F.4 "echo-test" performance comparison . . . . .		42
F.5 Enabling SSH . . . . .		43
F.6 Setting up a static IP address for PetaLinux . . . . .		43
<b>G References</b>		<b>44</b>

# 1 Introduction & motivation

This guide will present how to setup and use Xilinx's KRIA boards, in particular for running ROS2 on a host Ubuntu system, as well as for deploying micro-ROS[1] as a firmware on the micro-controller part of this board's chip.

The use of this device in particular is interesting because of the presence of a SoC comprising both a general purpose ARM core (Cortex A53), capable of running a Linux distribution, as well as another ARM core, real-time enabled, capable to run a RTOS (Cortex R5F). The figure 1 below shows a schematic view of the overall system we are trying to archive.

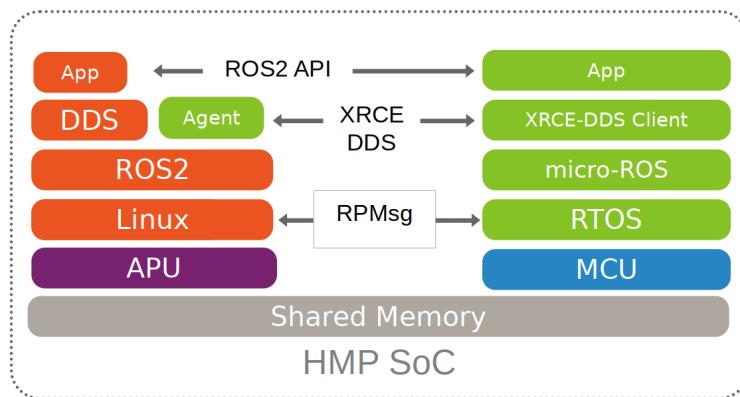


Figure 1: The Linux and ROS2 environment (orange) will communicate with the real-time, FreeRTOS and micro-ROS side (blue) using RPMsg (shared memory).

This document will give a step-by-step, chapter by chapter indication on how to go from a new board to a system ready to be used for testing internal communication for ROS, between both types of cores using the SoC shared memory.

## 2 Boot firmware update

The goal for the Linux side of the deployment is to have the latest LTS version of Ubuntu up and running.

In order to be able to boot such a newer version of Linux, the boot image of the board must first be updated.

The procedure is available in the official documentation<sup>1</sup>, but I will present it step by step here.

### 2.1 Getting the new firmware

A 2022.2 version of the board firmware is best in order to run the latest version of Ubuntu or PetaLinux properly.

The image download link can be obtained at the Atlassian page<sup>2</sup> on the topic, in the table detailing what version is best suited for which board.

In our case, we want to access this<sup>3</sup> page. AMD login will be asked, then download of firmware can be done.

### 2.2 Reaching the board recovery tool

Now the firmware .bin image is available<sup>4</sup>, it is possible to update it using the boards recovery tool. Here are the steps that must be taken in order to reach this tool and update the board:

- Power off the board and hold the firmware update button (FWUEN) when powering back the board.
- Connect the board to your machine via a Ethernet cable. This will obviously cut you internet access, so you should be set for that.
- In the case of the KR260, the bottom right port should be used, as seen in the figure 2 below.
- Select the wired network as your connection (must be "forced", since it doesn't have internet access). To do so, you should disable the IPv6, and set the IPv4 as manual, as visible in the figure 3 hereafter.
- Set a fixed IP address for your machine, in the 192.168.0.2/24 range, except the specific 192.168.0.111, which will be used by the board, and 192.168.0.1 which is the DNS and gateway. The netmask and gateway should also be respectively set to 255.255.255.0 and 192.168.0.1.
- Unplug the board power cord again. Hold the firmware update button (FWUEN) when powering back the board.
- Using a web browser on your host machine, access <http://192.168.0.111>. Thou shall now see the interface, as visible on the figure 4 below. If the page struggle to appear, you should try to un-plug and re-plug the Ethernet cable.

### 2.3 Updating the boot firmware

From this "recovery" page, it is possible to upload the .bin file downloaded previously onto the board using the "Recover Image" section at the bottom right of the page.

The board can be re-booted afterwards.

---

<sup>1</sup><https://docs.xilinx.com/r/en-US/ug1089-kv260-starter-kit/Firmware-Update>

<sup>2</sup><https://xilinx-wiki.atlassian.net/wiki/spaces/A/pages/1641152513/Kria+SOMs+Starter+Kits#K26-Boot-Firmware-Updates>

<sup>3</sup><https://www.xilinx.com/member/forms/download/design-license-xef.html?filename=BOOT-k26-starter-kit-20230516185703.bin>

<sup>4</sup>Something that looks like BOOT-k26-starter-kit-20230516185703.bin, at the time of writing this section.

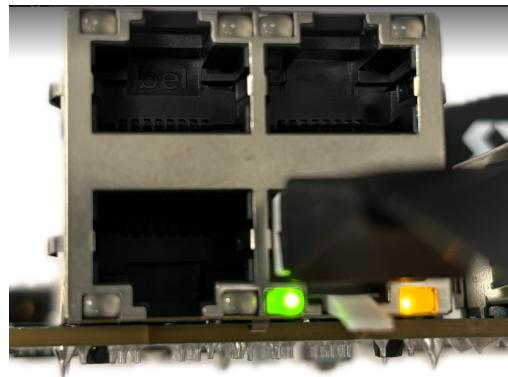


Figure 2: Port to be used for the Ethernet Boot Recovery Tool access on the KR260 board.

Cancel **Wired** Apply

Details Identity **IPv4** IPv6 Security

**IPv4 Method**

- Automatic (DHCP)
- Manual
- Shared to other computers

**Addresses**

Address	Netmask	Gateway
192.168.0.2	255.255.255.0	192.168.0.1

**DNS** Automatic

192.168.0.1

Separate IP addresses with commas

Figure 3: IPv4 settings for accessing the Recovery Tool.

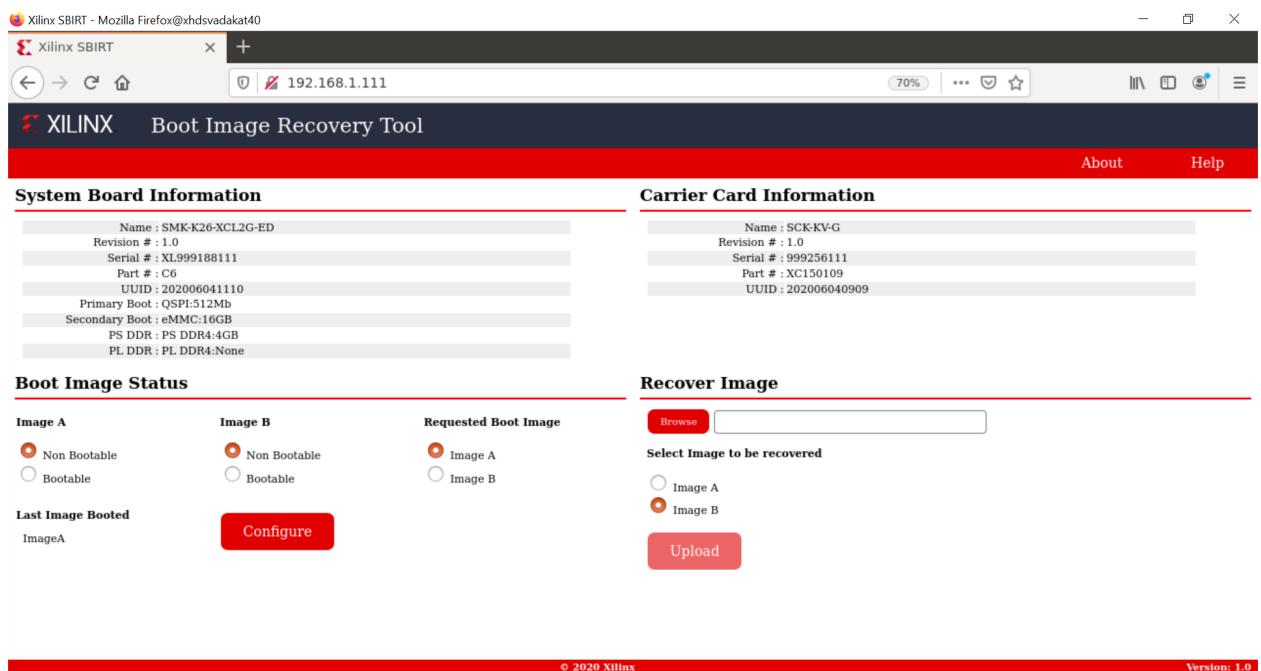


Figure 4: The recovery tool for the board, access from Firefox. We can see board information at the center, and the tools to upload the firmware at the bottom of the page.

### 3 Installing Ubuntu LTS 22.04

With the boot firmware being up-to-date, we can proceed to install a Linux distribution on our Kria board. The step needed to archive a full installation of Ubuntu LTS 22.04 will be presented in this section<sup>5</sup>. The figure 5 below shows where this operating system sits in the general system we are implementing.

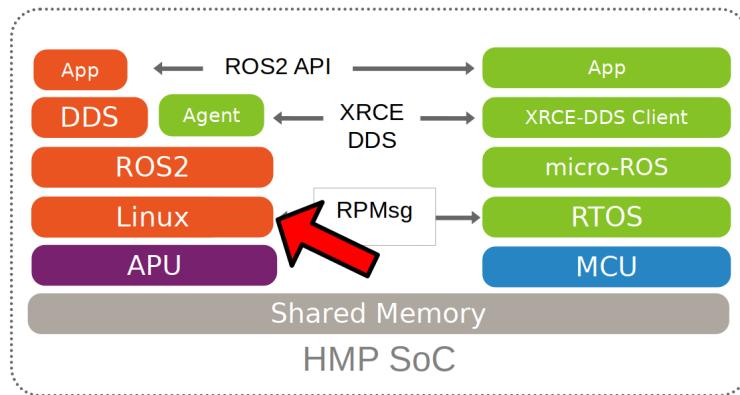


Figure 5: The Linux operating system (red border) runs on the APU (application, general purpose) side of the Kria board CPU. It is the base layer for the ROS2 system.

#### 3.1 Preparing and booting a Ubuntu 22.04 media

An official Ubuntu image exists and is provided by Xilinx, allowing the OS installation to be quick and straightforward. Ubuntu is a common and easy to use distribution. Furthermore, it allows to install ROS2 as a package, which is most convenient and will be done later in this guide.

Once the image has been downloaded at Canonical's page we can flash it onto the SD card, with the following instructions.

**DANGER:** The next part involve the dd command writing on disks!!! As always with the dd command, thou have to be **VERY** careful on what arguments thou give. Selecting the wrong disk will result on the destruction of thy data !! If you are unsure of what to do, seek assistance !

With the image available on thy machine and a SD card visible as /dev/sda device<sup>6</sup> one can simply run the dd command as follow to write the image to a previously formatted drive (here /dev/sda):

```
1 unxz iot-limerick-kria-classic-desktop-2204-20240304-165.img.xz
2 sudo dd if=iot-limerick-kria-classic-desktop-2204-20240304-165.img \
3     of=/dev/sda status=progress bs=8M && sync
```

Once the SD card is flashed and put back in the board, the micro-USB cable can be connected from the PC to the board. It is then possible to connect to the board in serial with an appropriate tool, for example picocom, as in the following example (the serial port that "appeared" was the /dev/ttyUSB1 in this case<sup>7</sup>, and the 115200 bit-rate is the default value for the board):

<sup>5</sup>The same procedure should work for other versions of Ubuntu, as long as they support the Kria board, but for this report and project, only the LTS 22.04 was tested (as of 2023-08-30).

<sup>6</sup>Again, it is critical to be 100\the correct device!

<sup>7</sup>If two boards are plugged in for serial, the second one will be /dev/ttyUSB5, then USB9 and so on.

```
1 sudo picocom /dev/ttyUSB1 -b 115200
```

In my case, I am using Emacs's serial-term:

```
1 M-x serial-term RET /dev/ttyUSB1 RET 115200 RET
```

The default username / password pair for the very first boot is `ubuntu` and `ubuntu`. You will then be prompted to enter a new password.

Once logged in, it is typically easier and more convenient to connect the board using SSH. When the board is connected to the network, it is possible to know its IP address with the `ip` command; then it is possible to connect to the board with `ssh`, as follow (example, with the first command to be run on the board and the second one on the host PC, both without the first placeholder hostnames):

```
1 kria# ip addr
2
3 host# ssh ubuntu@192.168.4.11
```

## 3.2 Network and admin setups

This section presents a variety of extra convenience configurations that can be used when setting-up the Kria board.

### 3.2.1 Static IP address

A static IP can be set by writing the following configuration into your `netplan` configuration file<sup>8</sup>.

The name of the files might vary:

```
1 sudo chmod 0600 /etc/netplan/50-cloud-init.yaml
2 sudo nano /etc/netplan/50-cloud-init.yaml
```

You can then set the wanted IP as follow<sup>9</sup>:

```
1
2   network:
3     renderer: NetworkManager
4     version: 2
5     ethernets:
6       eth0:
7         dhcp4: false
8         addresses:
9           - 192.168.11.107/24
10        routes:
11          - to: default
12            via: 192.168.11.1
13        nameservers:
14          addresses: [192.168.11.1]
```

Finally, the change in settings can be applied as follow:

```
1 sudo netplan apply
```

<sup>8</sup>The `chmod` command is used to update the permissions and silence some warnings

<sup>9</sup>For the routing part, it is key to have the `to` with a '-' in front of it; and then the `via` without, but aligned with the `t`.

### 3.2.2 [facultative] root password

**WARNING:** Depending on your use-case, the setup presented in this subsection can be a critical security breach as it removes the need for a root password to access the admin functions of the board's Linux. When in doubt, do not apply this configuration!!

If your board does not hold important data and is available to you only, for test or development, it might be convenient for the sudo tool to not ask for the password all the time. This change can be done by editing the sudoers file, and adding the parameter NOPASSWD at the sudo line:

```
1 sudo visudo  
2  
3 %sudo    ALL=(ALL:ALL) NOPASSWD: ALL
```

Again, this is merely a convenience setup for devices staying at your desk. If the board is meant to be used in any kind of production setup, a password should be set for making administration tasks.

With all of these settings, you should be able to update the software of your board without any issues:

```
1 sudo apt update  
2 sudo apt dist-upgrade  
3 sudo reboot now
```

### 3.2.3 Adding Xilinx specific repositories

The following commands will add PPA repositories that are specific for Xilinx boards using Ubuntu. It is then possible to update the package list and eventually upgrade to some new packages.

```
1 sudo add-apt-repository ppa:ubuntu-xilinx/updates  
2 sudo add-apt-repository ppa:xilinx-apps/ppa  
3 sudo apt update  
4 sudo apt upgrade
```

### 3.2.4 [facultative] Installing Docker

It is possible to have a version of Docker installed simply by using the available repository, but since we are on Ubuntu, a PPA is available from Docker in order to have the most up-to-date version.

Following the official documentation, the following steps can be taken to install the latest version of Docker on a Ubuntu system. The last command is meant to test the install. If everything went smoothly, you should see something similar to what is presented in the figure 6 below, after the commands:

```
1 sudo apt-get update  
2 sudo apt-get install ca-certificates curl  
3 sudo install -m 0755 -d /etc/apt/keyrings  
4 curl -fsSL https://download.docker.com/linux/ubuntu/gpg | \\\n      sudo gpg --dearmor -o /etc/apt/keyrings/docker.gpg  
6  
7 sudo chmod a+r /etc/apt/keyrings/docker.gpg  
8  
9 echo \\  
10   "deb [arch=$(dpkg --print-architecture)\" \\\n     signed-by=/etc/apt/keyrings/docker.gpg] \\\n     https://download.docker.com/linux/ubuntu \\\n     \"$(< . /etc/os-release && \\\n
```

```

14     echo "$VERSION_CODENAME)" stable" | \
15     sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
16
17 sudo apt-get update
18 sudo apt-get install docker-ce docker-ce-cli \
19   containerd.io docker-buildx-plugin docker-compose-plugin
20 sudo usermod -aG docker $USER
21 newgrp docker
22
23 docker run hello-world

```

```

ubuntu@kria:~$ docker run hello-world
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
70f5ac315c5a: Pull complete
Digest: sha256:fc6cf906cbfa013e80938cdf0bb199fbdbb86d6e3e013783e5a766f50f5dbce0
Status: Downloaded newer image for hello-world:latest

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
 1. The Docker client contacted the Docker daemon.
 2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
    (arm64v8)
 3. The Docker daemon created a new container from that image which runs the
    executable that produces the output you are currently reading.
 4. The Docker daemon streamed that output to the Docker client, which sent it
    to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
 $ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
 https://hub.docker.com/

For more examples and ideas, visit:
 https://docs.docker.com/get-started/

```

Figure 6: The return of a successful run of the hello world test Docker container.

### 3.2.5 [recommended] Purging snap

As the desktop-specific software are not used at all in the case of our project, there are some packages that can be purged in order for the system to become more lightweight.

In particular, the main issue with Ubuntu systems is the forced integration of Snap packages. Here are the command to use in order to remove all of that. These steps take a lot of time and need to be executed in that specific order<sup>10</sup>, but the system fan runs sensibly slower without all of this stuff.

Note that this all process is rather slow and can take up to 30min to complete.

```

1 sudo systemctl disable snapd.service
2 sudo systemctl disable snapd.socket
3 sudo systemctl disable snapd.seedied.service
4
5 sudo snap list #show installed package, remove them all:
6 sudo snap remove --purge firefox
7 sudo snap remove --purge gnome-3-38-2004
8 sudo snap remove --purge gnome-42-2204
9 sudo snap remove --purge gtk-common-themes

```

<sup>10</sup>The snap packages depends on each others. Dependencies cannot be remove before the package(s) that depends on them, thus the specific delete order.

```

10 sudo snap remove --purge snapd-desktop-integration
11 sudo snap remove --purge snap-store
12 sudo snap remove --purge bare
13 sudo snap remove --purge core20
14 sudo snap remove --purge core22
15 sudo snap remove --purge snapd
16 sudo snap list # check that everything is uninstalled
17
18 sudo rm -rf /var/cache/snapd/
19 sudo rm -rf ~/snap
20 sudo apt autoremove --purge snapd
21
22 # check once more that there is no more snap on the system
23 systemctl list-units | grep snapd

```

### 3.2.6 [facultative] Other unused heavy packages

Some other pieces of software can safely be removed since the desktop is not to be used.

```

1 sudo apt-get autoremove --purge yaru-theme-icon \
2   fonts-noto-cjk yaru-theme-gtk vim-runtime \
3   ubuntu-wallpapers-jammy humanity-icon-theme
4
5 sudo apt-get autoclean
6 sudo reboot now

```

### 3.2.7 [facultative] Slow boot services to disable

These packages (in particular the first one) are taking up a LOT of time at boot while providing no benefits<sup>11</sup>.

It is possible to disable them as follow:

```

1 sudo systemctl disable systemd-networkd-wait-online.service
2 sudo systemctl disable NetworkManager-wait-online.service
3 sudo systemctl disable cups.service
4 sudo systemctl disable docker.service
5 sudo systemctl disable containerd.service
6 sudo systemctl disable cloud-init-local.service

```

Additional, potentially unused services can be found using the very handy command:

```
1 sudo systemd-analyze blame
```

### 3.2.8 [facultative] Adding a swap partition

This part is very optional, in particular as it might slow down a bit the boot time of the board (~2s), however it might become handy to have swap memory available to avoid system failure under heavy use.

This whole procedure must be done externally, with the board system SD card mounted on a host PC as an external volume. As it is highly platform dependant, I will not give a detailed explanation on how to do it, yet here are the key points that should be done:

- Shutdown the Kria board, take out the SD card and put it in a host machine.
- Make sure the disk is visible.

---

<sup>11</sup>The CUPS and Docker services will be activated when used instead of during boot time.

- Make sure all volumes are **unmounted**.
- Resize the main root partition (**not** the boot) so a space the size of the wanted swap is free **after** the partition. You'd want something around 1GB.
- In the empty space, create a new partition, which type is "linux swap".
- Find and take note of the UUID of the new partition. This is useful hereafter.
- sync
- Un-mount everything, eject SD card.
- Put the SD card back in the Kria.
- Boot back to Ubuntu.

Going back on the Kria board Ubuntu after boot, the `/etc/fstab` file can be updated as follow, modulo your actual UUID for the newly created partition, to enable swap at boot time.

```

1 sudo -s
2 echo "UUID=8b13ed05-a91d-4x50-a44a-e654a0c67a2c none    swap    sw      0      0" >> /etc/fstab
3 reboot now

```

### 3.2.9 [OLD] Enabling remoteproc with Device-Tree patching

One of the advantage of this Kria board, as cited previously, is the presence of multiple types of core (APU, MCU, FPGA) on the same chip.

The part in focus in this guide is the usage of both the APU, running a Linux distribution and ROS2; and the MCU, running FreeRTOS and micro-ROS. Online available guides<sup>12</sup> ,<sup>13</sup> also provide information on how to deploy these types of systems and enabling remoteproc for the Kria board, but this guide will show a step-by-step, tried process to have a heterogeneous system up and running.

The communication between both side is meant to be done using shared memory, but some extra setup is required in order to be running the real-time firmware, in particular for deploying micro-ROS on it.

As a first step in that direction, this section of the report will present how to setup and use as an example firmware that utilizes the remoteproc device in Linux in order to access shared memory and communicate with the real-time firmware using the RPMsg system.

The communication system and interaction from the Linux side towards the real-time capable core is not enabled by default within the Ubuntu image provided by Xilinx.

In that regard, some modification of the device tree overlay (DTO) is required in order to have the remoteproc system starting.

- 1 Patching the device tree for RPMsg (standard, kernel space mode) Firstly, we need to get the original firmware device tree, converted into a readable format (DTS):

```

1 sudo dtc /sys/firmware/fdt 2> /dev/null > system.dts

```

Then, a custom-made patch file can be downloaded and applied. This file is available at the URL visible in the command below.

```

1 wget https://gitlab.com/sunoc/xilinx-kria-kv260-documentation/-/raw/main/src/system.patch
2
3 patch system.dts < system.patch

```

<sup>12</sup>A slideshow (JP) from Fixstar employees presents how to use the device tree to enable the communication between the cores.

<sup>13</sup>A blog post (JP) shows all major steps on how to enable the remoteproc.

2. Kernel cmd edit As for the board to be able to reserve the correct amount of memory with the new settings, some cma kernel configuration is needed<sup>14</sup>:

```
1 sudo nano /etc/default/flash-kernel
2
3 LINUX_KERNEL_CMDLINE="quiet splash cma=512M cpuidle.off=1"
4 LINUX_KERNEL_CMDLINE_DEFAULTS=""
5 sudo flash-kernel
```

Now the DTS file has been modified, one can regenerate the binary and place it on the /boot partition and reboot the board:

```
1 dtc -I dts -O dtb system.dts -o user-override.dtb
2 sudo mv user-override.dtb /boot/firmware/
3 sudo reboot now
```

3. Checking the patching After rebooting, you can check the content of the remoteproc system directory, and a remoteproc0 device should be visible, as follow:

```
1 ls /sys/class/remoteproc/
2 # remoteproc
```

If it is the case, it means that the patch was successful and that the remote processor is ready to be used!

---

<sup>14</sup>The overlapping memory will not prevent the board to boot, but it disables the PWM for the CPU fan, which will then run at full speed, making noise.

### 3.2.10 [NEW] Enabling remoteproc for RPMsg in userspace with device-tree patch

1. Kernel cmd edit

```
1 sudo nano /etc/default/flash-kernel
2
3 LINUX_KERNEL_CMDLINE="quiet splash cma=512M cpuidle.off=1"
4 LINUX_KERNEL_CMDLINE_DEFAULTS=""
5 sudo flash-kernel
```

2. Patching the device tree for RPMsg in userspace The system\_uio.patch

```
1 sudo dtc /sys/firmware/fdt 2> /dev/null > system.dts
2
3 wget https://gitlab.com/sunoc/xilinx-kria-kv260-documentation/-/raw/main/src/system_uio.patch
4
5 patch system.dts < system_uio.patch
6
7 dtc -I dts -O dtb system.dts -o user-override.dtb
8 sudo mv user-override.dtb /boot/firmware/
9 sudo reboot now
```

3. Checking the patching If everything went correctly, on reboot and re-logging in the device, you should see an output to the following command:

```
1 ls /sys/class/remoteproc/
```

### 3.2.11 [DEPRECATED] Proxy and DNS

**WARNING:** This setup is not needed anymore when using Honda-sensei's lab wired network. Adding it will cause the DNS to FAIL! This section is kept as a reference.

An issue that can occur when connecting the board to the internet is the conflicting situation with the university proxy. Indeed, as the network at Nanzan University requires to go through a proxy, some DNS errors appeared.

In that case, it might become needed to setup the proxy for the school.

This can be done as follow, by exporting a https base proxy configuration containing you AXIA credentials (this is specific to Nanzan University IT system), then by consolidating the configuration for other types of connections in the bashrc:

```
1 export https_proxy="http://<AXIA_username>:\n<AXIA_psw>@proxy.ic.nanzan-u.ac.jp:8080"
2
3
4 echo "export http_proxy=\"$https_proxy\" >> ~/.bashrc
5 echo "export https_proxy=\"$https_proxy\" >> ~/.bashrc
6 echo "export ftp_proxy=\"$https_proxy\" >> ~/.bashrc
7 echo "export no_proxy=\"localhost, 127.0.0.1,::1\" >> ~/.bashrc
```

Eventually the board can be rebooted in order for the setup to get applied cleanly.

### 3.2.12 [DEPRECATED] Jupyter notebook setup

Here are some instruction on how to install and setup Jupyter on a KRIA board, accessing it remotely and using it for making data analysis.

The following commands will set the required packages and install Jupyter itself<sup>15</sup>:

```
1 sudo apt-get update && sudo apt-get install python3 python3-pip python3-venv python3-virtualenv
2
3 virtualenv myjupyter
4 source ./myjupyter/bin/activate
5 python3 -m pip install jupyter pandas numpy matplotlib scipy
6
7 sudo reboot now
```

Then in a terminal on your host machine (not on the KRIA board), you can run the following command<sup>16</sup> to bind local ports:

```
1 ssh -L 8888:localhost:8888 ubuntu@192.168.11.107
```

Then on the opened SSH shell to the KRIA board:

```
1 source ./myjupyter/bin/activate
2 jupyter notebook
```

From there, it is possible to use the displayed URL (something that looks like `http://localhost:8888/tree?token`) to access the remote Notebook system from a local web browser. It is possible to do so with `localhost` since we have the `ssh` port map connection going on.

Eventually creating Notebooks and stuff, it is possible to obtain a situation like shown in the figure 7 below.

#### DDS measured data analysis

This notebook is meant to analyse the data collected from the measurement ROS2 node.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as stats

[3]: df = pd.read_csv('~/out.csv')
df = df.drop('Unnamed: 0', axis=1)
df = df.rename(columns={'16': '16 bytes RPMsg', '32': '32 bytes RPMsg'})

# Switchin to ms
df = df.mul(1000)

[5]: plt.rcParams['font.size']=10

scatterplot = df.plot(style='.', grid=True)
scatterplot.set_xlabel('Package sent [n]')
scatterplot.set_ylabel('Round trip time for one (1) byte [ms]')

[5]: Text(0, 0.5, 'Round trip time for one (1) byte [ms]')
```

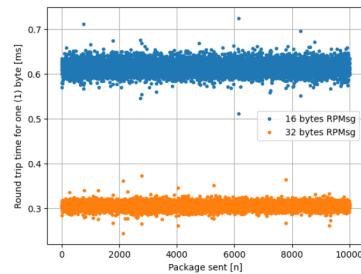


Figure 7: A test Jupyter Notebook for CSV data analysis.

<sup>15</sup>Alongside other packages useful for data analysis, such as `pandas` or `numpy`.

<sup>16</sup>In this example, the full `username@IP` is used, but a `.ssh/config` is also usable.

### 3.2.13 [TODO] Using a PetaLinux kernel in Ubuntu

TBD

## 3.3 RPMsg standalone evaluation

### 3.3.1 [TODO] RPMsg Cortex R5F demonstration firmware

### 3.3.2 RPMsg echo\_test software

In order to test the deployment of the firmware on the R5F side, and in particular to test the RPMsg function, we need some program on the Linux side of the KRIA board to "talk" with the real-time side.

Some source is provided by Xilinx to build a demonstration software that does this purpose: specifically interact with the demonstration firmware.

Here are the steps required to obtain the sources, and build the program.

As a reminder, this is meant to be done on the Linux running on the KRIA board, NOT on your host machine !

```
1 git clone https://github.com/Xilinx/meta-openamp.git
2 cd meta-openamp
3 git checkout xlnx-rel-v2022.2
4 cd ./recipes-openamp/rpmmsg-examples/rpmmsg-echo-test
5 make
6 sudo ln -s $(pwd)/echo_test /usr/bin/
```

Once this is done, it is possible to run the test program from the KRIA board's Ubuntu by running the echo\_test command.

## 4 Building micro-ROS as a static library

This section will present the way to build manually the micro-ROS system as static library that can be used to port it to a different platform.

### 4.1 Initial setup

In this section, the goal is to build the micro-ROS library in order to be able to integrate its functions into our Cortex R5F firmware.

All of this should be done via cross-compiling on a host machine, however it is most common in the guides about micro-ROS to build the firmware and libraries within a Docker, so we can have access of the ROS environment without installing it permanently.

One can simply run this command to summon a ROS2 Docker<sup>17</sup> with the wanted version, but first we also need to check the cross-compilation tools.

We are downloading the latest arm-none-eabi gcc compiler directly from the ARM website.

The cross-compilation tool can then be extracted, set as our toolchain variable, then passed as a parameter when creating the Docker container:

```
1 pushd /home/$USER/Downloads
2 wget
3   ↗ https://developer.arm.com/-/media/Files/downloads/gnu/13.2.rel1/binrel/arm-gnu-toolchain-13.2.rel1-x86_64-arm-none-eabi...
4 tar -xvf arm-gnu-toolchain-13.2.rel1-x86_64-arm-none-eabi.tar.xz
5 popd
6
7 toolchain="/home/$USER/Downloads/arm-gnu-toolchain-13.2.Rel1-x86_64-arm-none-eabi/"
```

<sup>17</sup>If Docker is not set up on your machine, you can follow the guide on the official website. When you can successfully run the "hello-world" container, you are good to go.

```
7
8
9 docker run -d --name ros_build -it --net=host \
10   --hostname ros_build \
11   -v $toolchain:/armr5-toolchain \
12   --privileged ros:iron
```

Now the container named `ros_build` was created, it is possible to "enter" it, and having access to the tools in it by running the following command that will open a bash shell in said container:

```
1 docker exec -it ros_build bash
```

## 4.2 Building the static library

Now we are in the ROS2 container, we can build the micro-ROS firmware as presented in the dedicated micro-ROS guide: #+BEGIN\_SRC sh echo 127.0.0.1 \$HOSTNAME >> /etc/hosts sudo apt update sudo apt-get -y install python3-pip \ wget \ nano  
. /opt/ros/\$ROS\_DISTRO/setup.bash  
mkdir microros\_ws cd microros\_ws git clone -b \$ROS\_DISTRO \ https://github.com/micro-ROS/micro\_ros \_setup.git \

## 5 ROS2 host system setup

This is focused on the ROS2 system being used on the Kria board as a base to run and / or build micro-ROS components.

The installation of ROS2 as a system will be presented, with two different ways of approaching the problem.

As for the previous section, the figure 8 below shows what part of the overall system we are talking about here.

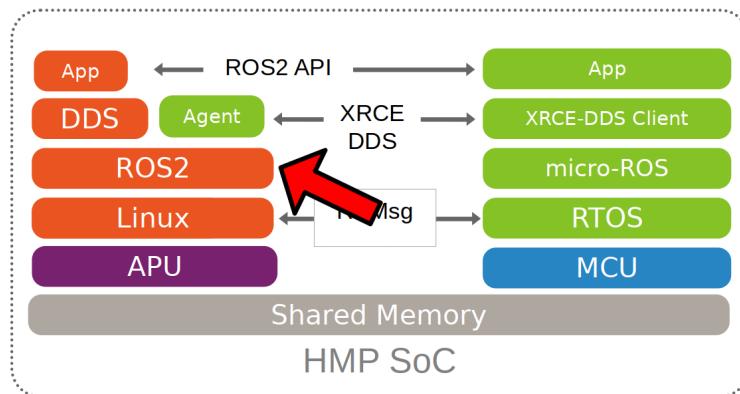


Figure 8: The ROS2 middle (red border) runs on top of the Linux, on the general-purpose core of the Kria board.

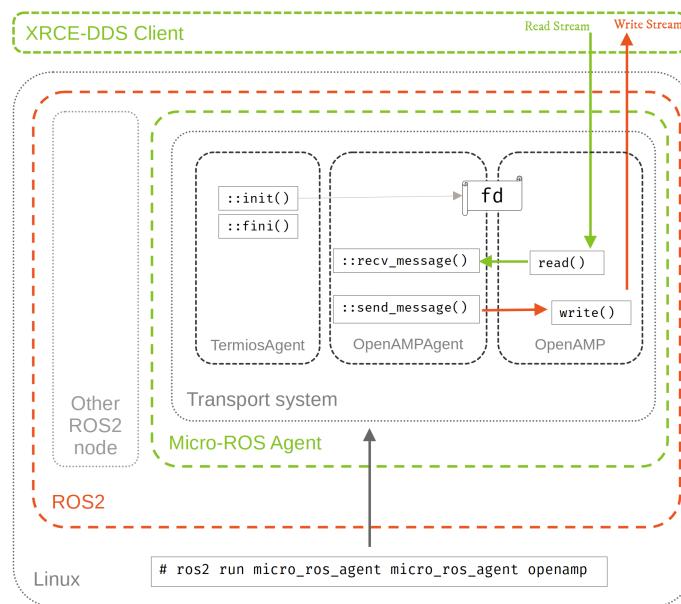


Figure 9: Methods architecture for the modified Agent node.

### 5.1 On the host Linux ("bare-metal")

Since an Ubuntu distribution is installed on the board, the installation of ROS2 can be done<sup>18</sup> in a standard way, using the repository.

<sup>18</sup>As always, this configuration was tested solely on Ubuntu LTS 22.04, with the ROS2 versions Humble and then Iron being deployed. Other combination of versions should work as well, but they are not tested for this guide. In case of doubt or problem, please refer to the official documentation.

An official documentation is provided with ROS2 themselves with a step-by-step guide on how to install ROS2 on a Ubuntu system<sup>19</sup>. We will be following this guide here<sup>19</sup>.

Firstly, we need to update the locals, enable the universe Ubuntu repository, get the key and add the repository for ROS2. This can be done as follow:

```
1 locale # check for UTF-8
2 sudo apt update && sudo apt install -y locales
3 sudo locale-gen en_US en_US.UTF-8
4 sudo update-locale LC_ALL=en_US.UTF-8 LANG=en_US.UTF-8
5 export LANG=en_US.UTF-8
6 locale # verify settings
7
8 sudo apt install -y software-properties-common
9 sudo add-apt-repository universe
10 sudo apt update && sudo apt install -y curl wget
11
12 wget https://raw.githubusercontent.com/ros/rosdistro/master/ros.key
13 sudo mv ros.key /usr/share/keyrings/ros-archive-keyring.gpg
```

Then, a thick one-liner is available to add the ROS2 repository to our system:

```
1 echo "deb [arch=$(dpkg --print-architecture) \
2 signed-by=/usr/share/keyrings/ros-archive-keyring.gpg] \
3 http://packages.ros.org/ros2/ubuntu $(. \
4 /etc/os-release && echo $UBUNTU_CODENAME) main" | \
5 sudo tee /etc/apt/sources.list.d/ros2.list > /dev/null
```

It is then possible to install ROS2<sup>20</sup> as follow:

```
1 sudo apt update
2 sudo apt upgrade -y
3 sudo apt install -y ros-$ROS_DISTRO-desktop \
4     ros-$ROS_DISTRO-ros-base \
5     python3-argcomplete \
6     ros-dev-tools
```

Once installed, it is possible to test the system with a provided example. You need to open two terminals and log wish SSH onto the board, then running respectively:

```
1 source /opt/ros/$ROS_DISTRO/setup.bash
2 ros2 run demo_nodes_cpp talker
```

And then:

```
1 source /opt/ros/$ROS_DISTRO/setup.bash
2 ros2 run demo_nodes_py listener
```

You should be able to see the first terminal sending "Hello world" messages, and the second one receiving them.

<sup>19</sup>The curl command from the guide does not work through the school proxy, but the command wget used instead does work. The key is then moved to the correct spot with mv.

<sup>20</sup>This command installs a complete "desktop" version of ROS2, containing many useful package for our project. If space is a constraint, different, less complete packages can be install. Please refer to the official documentation about it.

## 5.2 In a container (Docker)

As containers are used to test and build micro-ROS configurations, running ROS2 in a Docker is a great way to have a reproducible configuration of your system.

This part of the guide will present how to install Docker on the Kria board and then how to use it to deploy the latest version of ROS2. This section of the report assumes that Docker was installed on the target system.

The following commands will pull a ROS container, version iron, and name it `ros_build`.

A key part for having access to the interfaces (serial) is the mapping of the whole `/dev` range of devices from the host machine to the internal `/dev` of the container<sup>21</sup>. With the second command, we can execute bash as a way to open a terminal to the "inside" the container:

```
1 docker run -d --name ros_agent -it --net=host -v \
2     /dev:/dev --privileged ros:iron
3 docker exec -it ros_agent bash
```

From there, it becomes possible to simply use ROS2 as you would for a bare-metal install, and as presented in the section 5.1 above:

```
1 source /opt/ros/$ROS_DISTRO/setup.bash
2
3 # Create a workspace and download the micro-ROS tools
4 mkdir microros_ws
5 cd microros_ws
6 git clone -b $ROS_DISTRO https://github.com/micro-ROS/\
7     micro_ros_setup.git src/micro_ros_setup
8
9 # Update dependencies using rosdep
10 sudo apt update && rosdep update
11 rosdep install --from-paths src --ignore-src -y
12
13 # Install pip
14 sudo apt-get install python3-pip
15
16 # Build micro-ROS tools and source them
17 colcon build
18
19 # Download micro-ROS-Agent packages
20 source install/local_setup.bash
21 ros2 run micro_ros_setup create_agent_ws.sh
22
23 # Build step
24 ros2 run micro_ros_setup build_agent.sh
25
26 # Run a micro-ROS agent
27 ros2 run micro_ros_agent micro_ros_agent serial \
28     --dev /dev/ttyUSB1
```

Then once again in a similar way to the bare-metal deployment, it is possible to run a demonstration the ping-pong topic communication from a different shell<sup>22</sup>:

```
1 source /opt/ros/$ROS_DISTRO/setup.bash
2
```

<sup>21</sup>This is an example and this situation can become a security issue. It would be a better practice in a production environment to map only the devices that are actually in use.

<sup>22</sup>You need to be careful to have your shell in the "correct" space: these commands need to be run inside the container in which the previous setup were installed, not on the host running the container system. The hostname should help you to figure out where you are.

```
3 # Subscribe to micro-ROS ping topic  
4 ros2 topic echo /microROS/ping
```

## 6 Micro-ROS Client

## WORK\_IN\_PROGRESS

This is a Vitis IDE project meant to demonstrate the deployment of a micro-ROS system on the KRIA board, utilizing user-space RPMsg as a transport layer for communicating with an Agent on the Linux side.

### 6.1 Libraries

A goal of this branch is to try and keep the whole thing as portable as possible. Thus this repository ships will all the files for FreeRTOS and micro-ROS compilation, making it fairly heavy.

### 6.2 Prerequisite: KRIA board setup

This whole system is assumed to be run on a Xilinx KRIA KR260 or KV260 board. The setup of the Ubuntu setup, in particular the device tree configuration for the remote processor to work is needed.

This procedure can be found and followed on the dedicated guide available in my repo<sup>23</sup>.

### 6.3 Building the firmware

Apart for the libraries, an ARM cross-compiler is required for this system to be build. Here is how to proceed to set it up on a host system.

Download and extract a recent version of the arm-none-eabi compiler:

```
1 wget
→ https://developer.arm.com/-/media/Files/downloads/gnu/13.2.rel1/binrel/arm-gnu-toolchain-13.2.rel1-x86_64-arm-none-eabi...
2 tar -xvf arm-gnu-toolchain-13.2.rel1-x86_64-arm-none-eabi.tar.xz
3 rm -f arm-gnu-toolchain-13.2.rel1-x86_64-arm-none-eabi.tar.xz
```

Add some of the useful executable to the PATH with the correct name:

```
1 ln -sf /home/sunoc/Downloads/arm-gnu-toolchain-13.2.Rel1-x86_64-arm-none-eabi/bin/arm-none-eabi-gcc
→ /usr/bin/armr5-none-eabi-gcc
2 ln -sf /home/sunoc/Downloads/arm-gnu-toolchain-13.2.Rel1-x86_64-arm-none-eabi/bin/arm-none-eabi-size
→ /usr/bin/armr5-none-eabi-size
```

And then it is possible to use the provided Makefile to build the project:

```
1 make clean
2 make all
```

### 6.4 Deployment

The ROS Client firmware must be loaded from a compatible Linux distributing running on the APU of the same device.

Assuming you have a successfully built .elf file in the Debug directory, you can first send it to the target device<sup>24</sup> with a scp command, such as:

```
1 scp -O rpmsg_pingpong_microros_lib.elf pazuzu:~
```

<sup>23</sup>[https://gitlab.com/sunoc/xilinx-kria-kv260-documentation/-/blob/main/install\\_ubuntu.pdf?ref\\_type=heads](https://gitlab.com/sunoc/xilinx-kria-kv260-documentation/-/blob/main/install_ubuntu.pdf?ref_type=heads)

<sup>24</sup>In my case I have a SSH config for my board with the hostname pazuzu. Replace it with your own username@ip or a hostname config.

The following commands allows to remove the old .elf file, move the new one to the /lib/firmware dir and restart the board:

```
1 rm -f /lib/firmware/rpmsg_pingpong_microros_lib.elf
2 mv rpmsg_pingpong_microros_lib.elf /lib/firmware/
3 reboot now
```

If needed, the FPGA binary used for GPIO can be downloaded:

```
1 if [ ! -e bd_kr260_wrapper.bin ]
2 then
3     wget
4         ↳ https://gitlab.com/sunoc/saleae-logic-8-documentation/-/raw/7beca3a8fd2ab12e67bef57a290e27c6f408c2d9/bin/bd_kr260_wr...
```

Then, while logged on the board as root, one can load the new firmware as follow:

```
1 echo rpmsg_pingpong_microros_lib.elf > /sys/class/remoteproc/remoteproc0/firmware
2 echo start > /sys/class/remoteproc/remoteproc0/state
```

Once the Client is loaded, the communication with with the Argent and thus with the rest of the ROS ecosystem can start happening.

## 6.5 Starting the Agent

To build the Agent, a whole different building process is required. You can find the code in my fork of the eProxima repo<sup>25</sup>, and the building note here: [../libmicroros\\_kv260/agent\\_userspace.org](#)

## 6.6 Starting the Talker

A nice-looking ping pong Talker node in C++ can be build and run using it's own repository<sup>26</sup>. The process is a bit long, but once everything is built, you can run the Talker as follow:

```
1 . /opt/ros/iron/setup.sh
2 . install/setup.sh
3 ros2 run cpp_uros_monitor monitor
```

## 6.7 Other useful functions and notes

### 6.7.1 "Stop and Reboot" button

```
1 echo stop > /sys/class/remoteproc/remoteproc0/state
2 reboot now
```

### 6.7.2 Building the micro-ROS static library

TBD

<sup>25</sup><https://github.com/sunoc/Micro-XRCE-DDS-Agent/tree/uio>

<sup>26</sup><https://gitlab.com/sunoc/cpp-package-for-micro-ros-monitoring/-/tree/develop>

### 6.7.3 RPMsg in userspace

Only a few addresses must be changed in order for the Client to talk with a micro-ROS Agent running in userspace:

- RING\_TX in rsc\_table.c: should match with the DT
- RING\_RX in rsc\_table.c: should match with the DT
- SHARED\_MEM\_PA in platform\_info.c: should match with the DT
- SHARED\_MEM\_SIZE in platform\_info.c: should match with the DT

## 7 micro-ROS XRCE-DDS Agent

## WORK\_IN\_PROGRESS

The micro-ROS Agent on the ROS2 side is the last piece of the puzzle needed to allow our DDS environment to use RPMsg as a mean of communication, as visible on the schematic of the figure 10 below. In particular, it will be useful to modify this agent in order to archive the full RPMsg communication for ROS2<sup>27</sup>. An official documentation exists, but it gives little to no detail on how to deploy such modified, custom transport setup. This part of the guide will focus on it.

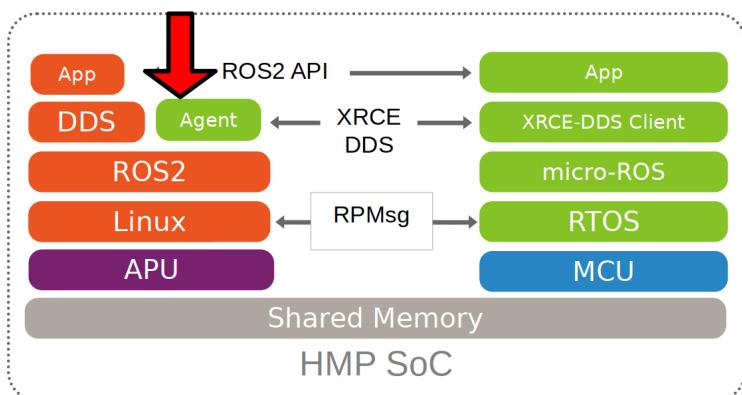


Figure 10: The agent (red border) allows for a micro-ROS instance to communicate with a ROS2 system. It is deployed on the Linux side, as a ROS2 node.

### 7.1 Building a XRCE-DDS agent in a Docker

The same command presented above for running a custom agent "bare-metal" can be run inside a Docker.

```
1 docker run -d --name XRCE_DDS_Agent -it --net=host -v \
2     /dev:/dev --privileged ros:iron
3     docker exec -it XRCE_DDS_Agent bash
```

```
1 git clone https://github.com/eProsima/Micro-XRCE-DDS-Agent.git
2 cd Micro-XRCE-DDS-Agent
3 docker build -t xrce-dds-agent .
4 docker run -it --privileged -v /dev:/dev xrce-dds-agent serial \
5     --dev /dev/ttyACM0
```

<sup>27</sup>This will be done with ROS2 agents "custom transport" system, which has little documentation. Some discussions about it exist though.

```
1 git clone https://github.com/eProsima/Micro-XRCE-DDS-Agent.git
2 cd Micro-XRCE-DDS-Agent
3 docker build -t xrce-dds-agent .
4 docker run -it --privileged -v /dev:/dev xrce-dds-agent serial \
5     --dev /dev/ttyACMO
```

## 8 Talker / ping-pong node

## WORK\_IN\_PROGRESS

A custom ping-pong node for ROS2 was developed in order to test the data transfer of the newly created Agent / Client system.

In order to have this node up and running, you need to set basically all the previously presented points:

- On the same Linux you plan to run the node, you will need to run and keep the Agent, as presented in the section 7.
- This ping-pong application is released as a ROS2 Python node. It is best to pull and run it in a ROS2 environment to avoid having dependencies issues, as presented in the section about ROS2.

## **9 Monitoring and performances evaluation WORK\_IN\_PROGRESS**

## **10 Conclusion & future**

**WORK\_IN\_PROGRESS**

## **11 Publications**

## **WORK\_IN\_PROGRESS**

Tryna link the conference and journal paper related to this project here. Unclear how atm.

## A DTO patch

This file is available in this repository: system.patch

```
1 --- system.dts.orig.2      2024-04-25 06:05:49.300357866 +0000
2 +++ system.dts          2024-09-05 08:57:00.690107364 +0000
3 @@ -1968,10 +1968,34 @@
4         no-map;
5         reg = <0x0 0x3ed88000 0x0 0x100000>;
6     };
7 - rproc_0_reserved: rproc_0_reserved@3ed00000 {
8 + rproc_0_reserved: rproc_0_reserved@3ec00000 {
9         no-map;
10        reg = <0x0 0x3ed00000 0x0 0x80000>;
11 +        reg = <0x0 0x3ec00000 0x0 0x180000>;
12     };
13 +
14 + rpu1vdev0vrинг0: rpu1vdev0vrинг0@3ef80000 {
15 +     no-map;
16 +     reg = <0x0 0x3ef80000 0x0 0x4000>;
17 + };
18 + rpu1vdev0vrинг1: rpu1vdev0vrинг1@3ef84000 {
19 +     no-map;
20 +     reg = <0x0 0x3ef84000 0x0 0x4000>;
21 + };
22 + rpu1vdev0buffer: rpu1vdev0buffer@3ef88000 {
23 +     no-map;
24 +     reg = <0x0 0x3ef88000 0x0 0x100000>;
25 + };
26 + rproc_1_reserved: rproc_1_reserved@3ef00000 {
27 +     no-map;
28 +     reg = <0x0 0x3ef00000 0x0 0x80000>;
29 + };
30 + uros_buf0: uros_buf@40000000 {
31 +     compatible = "shared-dma-pool";
32 +     reusable;
33 +     reg = <0x8 0x40000000 0x0 0x00400000>;
34 +     label = "uros_buf0";
35 + };
36 +
37 };
38 tcm_0a: tcm_0a@ffe00000 {
39     no-map;
40 @@ -1987,6 +2011,21 @@
41         compatible = "mmio-sram";
42         power-domain = <&zynqmp_firmware 16>;
43     };
44 + tcm_1a: tcm_1a@ffe90000 {
45 +     no-map;
46 +     reg = <0x0 0xffe90000 0x0 0x10000>;
47 +     status = "okay";
48 +     compatible = "mmio-sram";
49 +     power-domain = <&zynqmp_firmware 17>;
50 + };
51 + tcm_1b: tcm_1b@ffeb0000 {
```

```

52 +         no-map;
53 +         reg = <0x0 0xffeb0000 0x0 0x10000>;
54 +         status = "okay";
55 +         compatible = "mmio-sram";
56 +         power-domain = <&zynqmp_firmware 18>;
57 +     };
58 +
59     rf5ss@ff9a0000 {
60         compatible = "xlnx,zynqmp-r5-remoteproc";
61         xlnx,cluster-mode = <1>;
62 @@ -2004,8 +2043,20 @@
63             power-domain = <&zynqmp_firmware 7>;
64             mboxes = <&ipi_mailbox_rpu0 0>, <&ipi_mailbox_rpu0 1>;
65             mbox-names = "tx", "rx";
66 +
67     r5f_1 {
68         compatible = "xilinx,r5f";
69         #address-cells = <2>;
70         #size-cells = <2>;
71         ranges;
72         sram = <&tcm_1a &tcm_1b>;
73         memory-region = <&rproc_1_reserved>, <&rpu1vdev0buffer>,
74             <&rpu1vdev0ring0>, <&rpu1vdev0ring1>;
75         power-domain = <&zynqmp_firmware 8>;
76         mboxes = <&ipi_mailbox_rpu1 0>, <&ipi_mailbox_rpu1 1>;
77         mbox-names = "tx", "rx";
78     };
79 -
80     };
81 +
82     zynqmp_ipi1 {
83         compatible = "xlnx,zynqmp-ipi-mailbox";
84         interrupt-parent = <&gic>;
85 @@ -2028,6 +2079,46 @@
86             xlnx,ipi-id = <1>;
87         };
88     zynqmp_ipi2 {
89         compatible = "xlnx,zynqmp-ipi-mailbox";
90         interrupt-parent = <&gic>;
91         interrupts = <0 30 4>;
92         xlnx,ipi-id = <8>;
93         #address-cells = <1>;
94         #size-cells = <1>;
95         ranges;
96         /* APU<->RPU0 IPI mailbox controller */
97         ipi_mailbox_rpu1: mailbox@ff3f0b00 {
98             reg = <0xff3f0b00 0x20>,
99                 <0xff3f0b20 0x20>,
100                <0xff3f0940 0x20>,
101                <0xff3f0960 0x20>;
102             reg-names = "local_request_region",
103                 "local_response_region",
104                 "remote_request_region",

```

```

105      +           "remote_response_region";
106      +           # mbox-cells = <1>;
107      +           xlnx,ipi-id = <2>;
108      +       };
109      +   };
110      +   amba_pl: amba_pl@0 {
111      +       #address-cells = <2>;
112      +       #size-cells = <2>;
113      +       compatible = "simple-bus";
114      +       ranges ;
115      +       udmabuf@0x00 {
116      +           compatible = "ikwzm,u-dma-buf";
117      +           device-name = "udmabuf0";
118      +           size = <0x0 0x00200000>;
119      +           memory-region = <&uros_buf0>;
120      +       };
121      +       udmabuf@0x01 {
122      +           compatible = "ikwzm,u-dma-buf";
123      +           device-name = "udmabuf1";
124      +           size = <0x0 0x00200000>;
125      +           memory-region = <&uros_buf0>;
126      +       };
127      +   };
128
129      __symbols__ {
130      +         pinctrl_usb1_default =
131      +             "/firmware/zynqmp-firmware/pinctrl/usb1-default";

```

## B Custom tool-chain CMake settings

This file is available in this repository: custom\_r5f\_toolchain.cmake

```
1 set(CMAKE_SYSTEM_NAME Generic)
2 set(CMAKE_CROSSCOMPILING 1)
3 set(CMAKE_TRY_COMPILE_TARGET_TYPE STATIC_LIBRARY)
4
5
6 set(ARCH_CPU_FLAGS "-O2 -ffunction-sections -fdata-sections -fno-exceptions
7   ↳ -mcpu=cortex-r5 -mfpu=vfpv3-d16 -mfloat-abi=hard -Wl,--gc-sections -nostdlib --param
8     ↳ max-inline-insns-single=500 -DF_CPU=84000000L
9     ↳ -D'RCUTILS_LOG_MIN_SEVERITY=RCUTILS_LOG_MIN_SEVERITY_NONE')
10 set(ARCH_OPT_FLAGS "")
11
12 set(CMAKE_C_COMPILER arm-none-eabi-gcc)
13 set(CMAKE_CXX_COMPILER arm-none-eabi-g++)
14
15
16
17 set(__BIG_ENDIAN__ 0)
```

## C Custom Colcon meta settings

This file is available in this repository: custom r5f colcon.meta

```
1  {
2      "names": {
3          "tracetools": {
4              "cmake-args": [
5                  "-DTRACETOOLS_DISABLED=ON",
6                  "-DTRACETOOLS_STATUS_CHECKING_TOOL=OFF"
7              ]
8          },
9          "rosidl_typesupport": {
10             "cmake-args": [
11                 "-DROSIDL_TYPESUPPORT_SINGLE_TYPESUPPORT=ON"
12             ]
13         },
14         "rcl": {
15             "cmake-args": [
16                 "-DBUILD_TESTING=OFF",
17                 "-DRCL_COMMAND_LINE_ENABLED=OFF",
18                 "-DRCL_LOGGING_ENABLED=OFF"
19             ]
20         },
21         "rcutils": {
22             "cmake-args": [
23                 "-DENABLE_TESTING=OFF",
24                 "-DRCUTILS_NO_FILESYSTEM=ON",
25                 "-DRCUTILS_NO_THREAD_SUPPORT=ON",
26                 "-DRCUTILS_NO_64_ATOMIC=ON",
27                 "-DRCUTILS_AVOID_DYNAMIC_ALLOCATION=ON"
28             ]
29         },
30         "microxrcedds_client": {
31             "cmake-args": [
32                 "-DUCLIENT_PROFILE_UDP=OFF",
33                 "-DUCLIENT_PROFILE_TCP=OFF",
34                 "-DUCLIENT_PROFILE_DISCOVERY=OFF",
35                 "-DUCLIENT_PROFILE_SERIAL=OFF",
36                 "-UCLIENT_PROFILE_STREAM_FRAMING=OFF",
37                 "-DUCLIENT_PROFILE_CUSTOM_TRANSPORT=ON",
38                 "-DUCLIENT_CUSTOM_TRANSPORT_MTU=8232"
39             ]
40         },
41     },
42     "rmw_microxrcedds": {
43         "cmake-args": [
44             "-DRMW_UXRCE_MAX_NODES=1",
45             "-DRMW_UXRCE_MAX_PUBLISHERS=5",
46             "-DRMW_UXRCE_MAX_SUBSCRIPTIONS=5",
47             "-DRMW_UXRCE_MAX_SERVICES=1",
48             "-DRMW_UXRCE_MAX_CLIENTS=1",
49             "-DRMW_UXRCE_MAX_HISTORY=4",
50             "-DRMW_UXRCE_TRANSPORT=custom"
51         ]
52     }
53 }
```

52              }  
53        }  
54    }

## D Firmware time functions

### D.1 main

This file is available in this repository: `clock.c` but a potentially more up-to-date version is visible directly at the `libmicroros_kv260` repository: `clock.c`

```
1 #include "microros.h"
2
3
4 int _gettimeofday( struct timeval *tv, void *tzvp )
5 {
6     XTime t = 0;
7     XTime_GetTime(&t); //get uptime in nanoseconds
8     tv->tv_sec = t / 1000000000; // convert to seconds
9     tv->tv_usec = ( t % 1000000000 ) / 1000; // get remaining microseconds
10    return 0; // return non-zero for error
11 } // end _gettimeofday()
12
13
14 void UTILS_NanosecondsToTimespec( int64_t llSource,
15                                     struct timespec * const pxDestination )
16 {
17     long lCarrySec = 0;
18
19     /* Convert to timespec. */
20     pxDestination->tv_sec = ( time_t ) ( llSource / NANOSECONDS_PER_SECOND );
21     pxDestination->tv_nsec = ( long ) ( llSource % NANOSECONDS_PER_SECOND );
22
23     /* Subtract from tv_sec if tv_nsec < 0. */
24     if( pxDestination->tv_nsec < 0L )
25     {
26         /* Compute the number of seconds to carry. */
27         lCarrySec = ( pxDestination->tv_nsec / ( long ) NANOSECONDS_PER_SECOND ) + 1L;
28
29         pxDestination->tv_sec -= ( time_t ) ( lCarrySec );
30         pxDestination->tv_nsec += lCarrySec * ( long ) NANOSECONDS_PER_SECOND;
31     }
32 }
33
34 int clock_gettime( clockid_t clock_id,
35                     struct timespec * tp )
36 {
37     TimeOut_t xCurrentTime = { 0 };
38
39     /* Intermediate variable used to convert TimeOut_t to struct timespec.
40      * Also used to detect overflow issues. It must be unsigned because the
41      * behavior of signed integer overflow is undefined. */
42     uint64_t ullTickCount = OULL;
43
44     /* Silence warnings about unused parameters. */
45     ( void ) clock_id;
46
47     /* Get the current tick count and overflow count. vTaskSetTimeOutState()
48      * is used to get these values because they are both static in tasks.c. */
```

```

49     vTaskSetTimeOutState( &xCurrentTime );
50
51     /* Adjust the tick count for the number of times a TickType_t has overflowed.
52      * portMAX_DELAY should be the maximum value of a TickType_t. */
53     ullTickCount = ( uint64_t ) ( xCurrentTime.xOverflowCount ) << ( sizeof( TickType_t
54     → ) * 8 );
55
56     /* Add the current tick count. */
57     ullTickCount += xCurrentTime.xTimeOnEntering;
58
59     /* Convert ullTickCount to timespec. */
60     UTILS_NanosecondsToTimespec( ( int64_t ) ullTickCount * NANOSECONDS_PER_TICK, tp );
61
62     return 0;
63 }
```

## D.2 header file

```

1  /**< Microseconds per second. */
2  #define MICROSECONDS_PER_SECOND    ( 1000000LL )
3  /**< Nanoseconds per second. */
4  #define NANOSECONDS_PER_SECOND     ( 1000000000LL )
5  /**< Nanoseconds per FreeRTOS tick. */
6  #define NANOSECONDS_PER_TICK       ( NANOSECONDS_PER_SECOND / configTICK_RATE_HZ )
```

## E Firmware memory allocation functions

### E.1 main

This file is available in this repository: allocators.c but a potentially more up-to-date version is visible directly at the libmicroros\_kv260 repository: allocators.c

```
1 #include "allocators.h"
2
3 //int absoluteUsedMemory = 0;
4 //int usedMemory = 0;
5
6 void * __freertos_allocate(size_t size, void * state){
7     (void) state;
8     LPRINTF("-- Alloc %d (prev: %d B)\r\n", size, xPortGetFreeHeapSize());
9     // absoluteUsedMemory += size;
10    // usedMemory += size;
11
12    LPRINTF("Return for the allocate function w parameter size = %d\r\n", size);
13
14    return pvPortMalloc(size);
15}
16
17 void __freertos_deallocate(void * pointer, void * state){
18     (void) state;
19     LPRINTF("-- Free 0x%x (prev: %d B)\r\n", pointer, xPortGetFreeHeapSize());
20     if (NULL != pointer)
21     {
22         LPRINTF("Pointer is not null.\r\n");
23         // usedMemory -= getBlockSize(pointer);
24         // LPRINTF("usedMemory var updated: %d\r\n", usedMemory);
25         vPortFree(pointer);
26     }
27     else
28     {
29         LPERROR("Trying to deallocate a null pointed. Doing nothing.\r\n");
30     }
31}
32
33 void * __freertos_reallocate(void * pointer, size_t size, void * state){
34     (void) state;
35     LPRINTF("--- Realloc 0x%x -> %d (prev: %d B)\r\n", pointer, size,
36     ↵ xPortGetFreeHeapSize());
37     // absoluteUsedMemory += size;
38     // usedMemory += size;
39     if (NULL == pointer)
40     {
41         return __freertos_allocate(size, state);
42     }
43     else
44     {
45         // usedMemory -= getBlockSize(pointer);
46         // return pvPortRealloc(pointer,size);
47 }
```

```

48     __freertos_deallocate(pointer, state);
49     return __freertos_allocate(size, state);
50 }
51 }
52 }
53
54 void * __freertos_zero_allocate(size_t number_of_elements, size_t size_of_element, void
55     * state){
56     (void) state;
57     LPRINTF("<!-- Calloc %d x %d = %d -&gt; (prev: %d
58         B)\r\n", number_of_elements, size_of_element, number_of_elements * size_of_element,
59         xPortGetFreeHeapSize());
60 //    absoluteUsedMemory += number_of_elements * size_of_element;
61 //    usedMemory += number_of_elements * size_of_element;
62
63     return pvPortCalloc(number_of_elements, size_of_element);
64 }</pre>

```

## E.2 header file

```

1 #ifndef _ALLOCATORS_H_
2 #define _ALLOCATORS_H_
3
4 #include "microros.h"
5
6 extern int absoluteUsedMemory;
7 extern int usedMemory;
8
9
10 void * __freertos_allocate(size_t size, void * state);
11 void __freertos_deallocate(void * pointer, void * state);
12 void * __freertos_reallocate(void * pointer, size_t size, void * state);
13 void * __freertos_zero_allocate(size_t number_of_elements,
14                               size_t size_of_element, void * state);
15
16 #endif // _ALLOCATORS_H_
```

## F [deprecated] Installing Linux (PetaLinux option)

This part is an alternative to the previous section 3, where the installation of PetaLinux as a host system<sup>28</sup> will be presented, as being an alternative to the Ubuntu LTS. Here are some reasoning why you'd prefer to use PetaLinux instead of Ubuntu:

- Direct support from Xilinx. While Ubuntu is also an official port, it is provided through Canonical.
- If you plan to use ROS2 in containers.

It is also to be noted that the building process of a PetaLinux image is not as trivial as simply flashing a downloaded ISO file. It takes some time to configure it and even more time to be built, especially if you are working on a lower-end machine.

<sup>28</sup>It is to be noted that for this part, the "other" KRIA board (KR260) was used instead of the KV260 that was being used for everything tried on top of Ubuntu. No significant difference should be noted.

## F.1 General installation

The indication on how to install PetaLinux on a KRIA board is detailed in the official documentation<sup>29</sup>, but for the KV260 board with an older release of the Linux. This guide will try to keep an up-to-date version of the guide, in particular for the KR260 KRIA board and, at the time of writing, the PetaLinux in version 2023.2 for the installer. The 2022.2 version is recommended for this board's latest updated boot firmware, as presented in the Section Boot firmware above; however I happened to have issues with the building process from a system based on Ubuntu 22.04. Thus the choice of the 2023.2 that should be compatible.

Two tools should be downloaded upfront in order to be ready and create the PetaLinux image for our target board. Both can be obtained at the Xilinx official download page<sup>30</sup>. Login with an AMD account is required in order to download the files we need, namely, we'll have to get:

- The installer tool: `petalinux-v2023.2-10121855-installer.run` or similar name.
- The SOM board support package: `xilinx-kr260-starterkit-v2023.2-10140544.bsp` or similar name<sup>31</sup>.

The specific name of the utils you'll get to download will evolve over time.

For the following command, I will assume that both files were downloaded into the same directory and that you have a shell open in said directory.

Firstly, a crap ton of dependencies are needed. They are detailed in the release notes of the version of the tool you are downloading and it depends on your distribution, but here is a one-liner that works for the PetaLinux 2023.2 to be built on Ubuntu(-based) 22.04 LTS:

```
1 sudo apt-get install iproute2 gawk python3 build-essential gcc git make \
2     net-tools libncurses5-dev tftpd zlib1g-dev libssl-dev flex bison libselinux1 \
3     gnupg wget git-core diffstat chrpath socat xterm autoconf libtool tar unzip \
4     texinfo zlib1g-dev gcc-multilib automake zlib1g:i386 screen pax gzip cpio \
5     python3-pip python3-pexpect xz-utils debianutils iputils-ping python3-git \
6     python3-jinja2 libegl1-mesa libsdl1.2-dev pylint asciidoc
```

Testing, if building on Debian 12 the following variant of the same command has been confirmed to work:

```
1 sudo apt-get install iproute2 gawk python3 build-essential gcc git make \
2     net-tools libncurses5-dev tftpd-hpa zlib1g-dev libssl-dev flex bison libselinux1 \
3     gnupg wget git-core diffstat chrpath socat xterm autoconf libtool tar unzip \
4     texinfo zlib1g-dev gcc-multilib automake zlib1g screen pax gzip cpio python3-pip \
5     python3-pexpect xz-utils debianutils iputils-ping python3-git python3-jinja2 \
6     libegl1-mesa libsdl1.2-dev pylint asciidoc
```

Installing the petalinux tools can be done with the `.run` script<sup>33</sup>, as follow:

```
1 sudo chmod +x ./petalinux-v2023.2-10121855-installer.run
2 ./petalinux-v2023.2-10121855-installer.run
3 bash
4 source settings.sh
```

A project can then be created, using the `.bsp` file. From there it is possible to build the image file meant to be booted by the board:

<sup>29</sup>[https://xilinx.github.io/kria-apps-docs/kv260/2021.1/build/html/docs/build\\_petalinux.html](https://xilinx.github.io/kria-apps-docs/kv260/2021.1/build/html/docs/build_petalinux.html)

<sup>30</sup><https://www.xilinx.com/support/download/index.html/content/xilinx/en/downloadNav/embedded-design-tools/2022-2.html>

<sup>31</sup>As of the time of writing this section, it was recommended<sup>32</sup>to be using the boot firmware in version 2022.2. Thus the software pack for the board will also be using this version.

<sup>32</sup><https://xilinx-wiki.atlassian.net/wiki/spaces/A/pages/1641152513/Kria+SOMs+Starter+Kits#K26-Boot-Firmware-Updates>

<sup>33</sup>It is to be noted that the `source` command will run best on bash. Trying to run it on zsh for ex. will cause issues. This can be enforced with the command `sudo dpkg-reconfigure bash`.

```

1 petalinux-create -t project -s xilinx-kr260-starterkit-v2023.2-10140544.bsp
2 cd xilinx-kr260-starterkit-2023.2/

```

The more general option allows to enable stuff related to hardware configuration<sup>34</sup>. For this step, it is possible to select the openAMP setting the the DTG Settings, as visible in the Figure 11 below:

```

1 petalinux-config --get-hw-description ./hardware/xilinx-kr260-starterkit-2023.2/kr260_starter_kit.xsa

```

Then, here are the options that we want to enable:

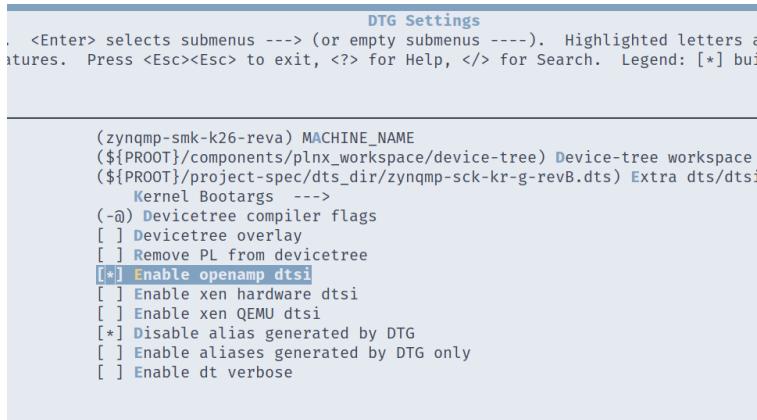


Figure 11: In the hardware config, it is possible to select the openamp device tree source.

Some kernel and module configuration are needed in order for the petalinux system to work properly. These setup can be set using the following command:

```

1 petalinux-config -c rootfs

```

Then, here are the options that we want to enable:

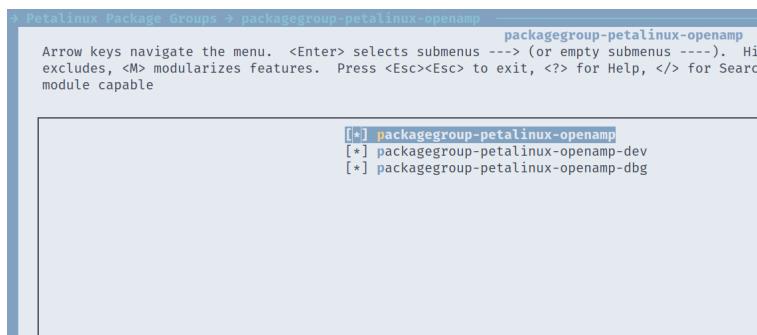


Figure 12: In the rootfs config, it is possible to select the openamp package.

It is then possible to build<sup>35</sup> and package the image. Most critically compared to the example, the openamp.dtb overlay must be used in order for OpenAMP and subsequently RPMMsg to become usable.

<sup>34</sup>If you search in the work directory, you might notice that three differently named .xsa files exist. After a quick diff check, they appear to be exactly identical.

<sup>35</sup>Beware, the petalinux-build command execution takes a lot of time.

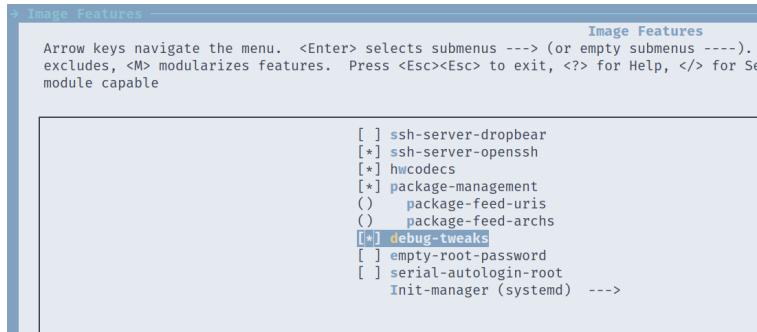


Figure 13: OpenSSH can be configured in the same rootfs config.

```

1 petalinux-build
2 petalinux-package --wic --bootfiles "ramdisk.cpio.gz.u-boot boot.scr Image system.dtb"

```

**DANGER:** The next part involve the dd command writing on disks!!! As always with the dd command, thou have to be **VERY** careful on what arguments thou give. Selecting the wrong disk will result on the destruction of thy data !! If you are unsure of what to do, seek assistance !

Finally, and now an image has been generated, it can be burnt to a micro-SD card the usual way:

```

1 sudo dd if=./images/linux/petalinux-sdimage.wic \
2   of=/dev/sda status=progress bs=8M && sync

```

From that point, the media can be ejected, put in the board and the boot process should happen successfully.

## F.2 First login

On the first start, if you have a serial terminal open to the KR260 board, you'll be prompted to enter a login. The default user is petalinux. You'll then be prompted to enter a password.

## F.3 Testing openamp "echo-test"

A good and out-of-the-box way to try and confirm that the OpenAMP setup was done properly is to run the so-called "echo-test". This program will basically send series of packages to the R5F core and expect them to be returned.

As presented in the official documentation<sup>36</sup>, this can be simply done with the following commands. If the return of the echo\_test commands states that zero error occurs, the system should be good to go:

```

1 sudo -s
2 echo image_echo_test > /sys/class/remoteproc/remoteproc0/firmware
3 echo start > /sys/class/remoteproc/remoteproc0/state
4 echo_test
5 echo stop > /sys/class/remoteproc/remoteproc0/state

```

<sup>36</sup>[https://xilinx.github.io/kria-apps-docs/openamp/build/html/openamp\\_landing.html#openamp-demos](https://xilinx.github.io/kria-apps-docs/openamp/build/html/openamp_landing.html#openamp-demos)

## F.4 "echo-test" performance comparison

Now we have a working demonstration for RPMsg, a goal was to compare this system between boards.

Here is the complete sequence in order to archive a timely measured echo\_test run, namely:

- Installation of the build dependencies
- Clone of the Xilinx repository
- Patching of the main source to add the time measurement
- Building of the test software
- Communication with the R5F, with reload of the firmware

```

1 sudo dnf install git make gcc binutils packagegroup-core-buildessential
2 sudo ln -s /usr/bin/aarch64-xilinx-linux-gcc /bin/cc
3 git clone https://github.com/OpenAMP/openamp-system-reference.git
4 cd openamp-system-reference/examples/linux/rpmsg-echo-test/
5 wget https://gitlab.com/sunoc/xilinx-kria-kv260-documentation/-/raw/main/src/echo_test.patch
6 patch -u -b echo_test.c -i echo_test.patch
7 make
8 sudo -s
9 echo image_echo_test > /sys/class/remoteproc/remoteproc0/firmware
10 echo start > /sys/class/remoteproc/remoteproc0/state
11 ./echo_test
12 echo stop > /sys/class/remoteproc/remoteproc0/state

```

The figure 14 below shows a comparison for the data throughput that is reachable for a packages round trip of individual RPMsg packages for both PetaLinux running on the KR260 board and Ubuntu running on the KV260.

A difference exist but it is not as significant as to explain the rather slow transfer rate when used in combination with ROS2 DDS.

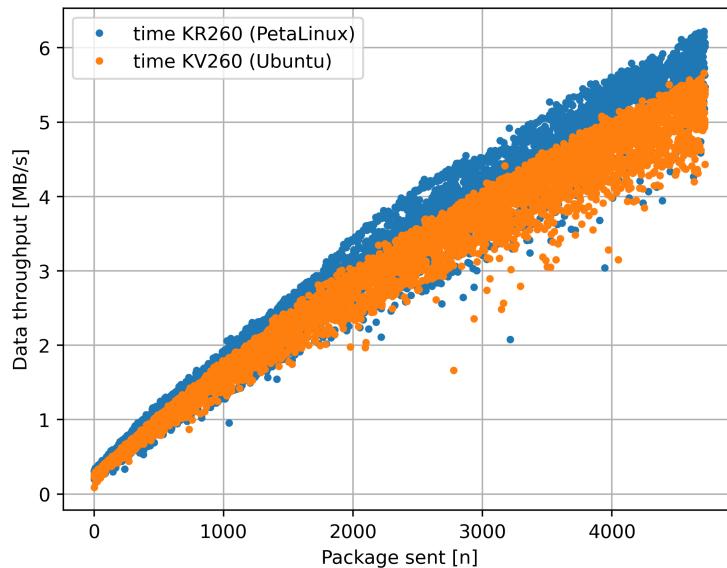


Figure 14: A graph of the data transfer rate for individual packages to be transmitted and received back from the Linux.

This fast delay for data transmission was later confirmed by measurement done using an external logic analyser<sup>37</sup>.

## F.5 Enabling SSH

## F.6 Setting up a static IP address for PetaLinux

This is not a trivial task in petalinux.

As it appears, the classic /etc/network/interfaces must be modified, however and for some unknown reason the DHCP keep being used even when disabled, and the actual static interface is not used until restarted.

Multiple steps are required in order to mitigate all of this.

First, we can make our standard interfaces configuration, in the /etc/network/interfaces:

```
1 # Wired or wireless interfaces
2 auto eth0
3 iface eth0 inet static
4   address 192.168.11.107
5   netmask 255.255.255.0
6   network 192.168.11.0
7   gateway 192.168.11.1
```

Then we need to have systemd to run a script after boot time to restart the eth0 interface.

A very basic script, as follow need to be placed somewhere on the system, for example here, in the /home/root/net.sh:

```
1#!/bin/sh -e
2ifdown eth0
3ifup eth0
4exit 0
```

Then, we want to create a systemd service that will run this script as idle, i.e. after all the rest. Said file, in our case named updown.service must be placed in the /etc/systemd/system/ directory:

```
1 [Unit]
2 Description=Restart once more the eth0 network to enable the static IP
3
4 [Service]
5 Type=idle
6 ExecStart=/bin/sh /home/root/net.sh
7
8 [Install]
9 WantedBy=multi-user.target
```

Finally, as a root we need to set the correct right for these files and finally enabling the service before reboot:

```
1 sudo -s
2 chmod +x /home/root/net.sh
3 chmod 644 /etc/systemd/system/updown.service
4 systemctl enable updown.service
5 reboot now
```

With all of this, on the next reboot of the board, the static IP should be available alongside the DHCP IP.

---

<sup>37</sup><https://gitlab.com/sunoc/saleae-logic-8-documentation>

## G References

### References

- [1] K. Belsare, A. C. Rodriguez, P. G. Sanchez, J. Hierro, T. Kolcon, R. Lange, I. Lutkebohle, A. Malki, J. M. Losa, F. Melendez, M. M. Rodriguez, A. Nordmann, J. Staschulat, and J. von Mendel, *Micro-ROS*. Cham: Springer International Publishing, 2023.